

# Neural Networks

## Part 1

Jon Dehdari

January 28, 2016

## Extending Logistic Regression (=Softmax Regression)


- Recall that **logistic regression** involves the dot product of an input vector and a weight matrix, then a normalized sigmoid function (softmax)

## Extending Logistic Regression (=Softmax Regression)

- Recall that **logistic regression** involves the dot product of an input vector and a weight matrix, then a normalized sigmoid function (softmax)
- A **feedforward neural network** just adds one or more layers between the input vector and the softmax output


# Extending Logistic Regression (=Softmax Regression)

- Recall that **logistic regression** involves the dot product of an input vector and a weight matrix, then a normalized sigmoid function (softmax)
- A **feedforward neural network** just adds one or more layers between the input vector and the softmax output









## Why Use Hidden Layers?

- In contrast to log-linear models, neural networks can have **non-linear** representations of data
- The **universal approximation theorem** (George Cybenko, 1989) found that a neural network with one hidden layer can approximate **any continuous function**
- A network with two hidden layers can represent discontinuous functions





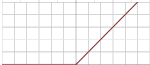



## Activation Functions ( $\sigma$ )

In each layer, the output of the dot product goes through an **activation function** ( $\sigma$ ). Here are some examples:

Name	Visualization	$f(x) =$	Notes
Linear (= Identity)		$x$	Not useful for hidden layers
Heaviside Step		$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	Not differentiable
Rectified Linear (ReLU)		$\begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	Surprisingly useful in practice
Tanh		$\frac{2}{1+e^{-2x}} - 1$	A soft step function; ranges from -1 to 1
Logistic ('sigmoid')		$\frac{1}{1+e^{-x}}$	Another soft step function; ranges from 0 to 1
Softmax		$\frac{e^{W_y \cdot x}}{Z}$	Normalized sigmoidal function. Useful for last layer when training on cross entropy

## Activation Functions ( $\sigma$ )

In each layer, the output of the dot product goes through an **activation function** ( $\sigma$ ). Here are some examples:

Name	Visualization	$f(x) =$	Notes
Linear (= Identity)		$x$	Not useful for hidden layers
Heaviside Step		$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	Not differentiable
Rectified Linear (ReLU)		$\begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	Surprisingly useful in practice
Tanh		$\frac{2}{1+e^{-2x}} - 1$	A soft step function; ranges from -1 to 1
Logistic ('sigmoid')		$\frac{1}{1+e^{-x}}$	Another soft step function; ranges from 0 to 1
Softmax		$\frac{e^{W_y \cdot x}}{Z}$	Normalized sigmoidal function. Useful for last layer when training on cross entropy

List of activation functions in Keras: [keras.io/activations](https://keras.io/activations)

# Training Neural Networks

- At a high level, the weights in a neural net are set by means of the blame game – whenever it guesses incorrectly, change the weights that were the most responsible for making that guess



# Training Neural Networks

- At a high level, the weights in a neural net are set by means of the blame game – whenever it guesses incorrectly, change the weights that were the most responsible for making that guess
- Whenever the network guesses a training instance correctly, don't change anything

# Training Neural Networks

- At a high level, the weights in a neural net are set by means of the blame game – whenever it guesses incorrectly, change the weights that were the most responsible for making that guess
- Whenever the network guesses a training instance correctly, don't change anything
- The weights are usually trained by a form of the gradient descent optimization algorithm
- The gradients are calculated by error **backpropagation**
- First, do a normal forward pass through the network, to determine the **error/loss** (how different the output was from the 'correct' answer)
- Then, do a backwards pass (end to start), changing the weights to minimize errors

# Loss / Objective Functions

- **Discrete Outputs:**

- Binary Cross-Entropy (0-1 loss): 0 if correct, 1 if incorrect
- Categorical Cross-Entropy: good old cross-entropy. Eg.
  - 0 if  $p(y) = 1.0$ ,
  - 1 if  $p(y) = 0.5$ ,
  - 2 if  $p(y) = 0.25$ ,
  - 3 if  $p(y) = 0.125$ ,
  - ...

- **Continuous Outputs:**

- Mean Squared Error (MSE):  $\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$
- Root Mean Squared Error (RMSE):  $\sqrt{MSE}$
- Mean Absolute Error (MAE):  $\frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$

## Loss / Objective Functions

- **Discrete Outputs:**

- Binary Cross-Entropy (0-1 loss): 0 if correct, 1 if incorrect
- Categorical Cross-Entropy: good old cross-entropy. Eg.
  - 0 if  $p(y) = 1.0$ ,
  - 1 if  $p(y) = 0.5$ ,
  - 2 if  $p(y) = 0.25$ ,
  - 3 if  $p(y) = 0.125$ ,
  - ...

- **Continuous Outputs:**

- Mean Squared Error (MSE):  $\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$
- Root Mean Squared Error (RMSE):  $\sqrt{MSE}$
- Mean Absolute Error (MAE):  $\frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$

List of loss functions in Keras: [keras.io/objectives](https://keras.io/objectives)

## Autoencoders

- An **autoencoder** is a neural network where the size of the output layer is the same size as the input layer
- The hidden layers are usually smaller
- The goal is to generalize the training data
- Since no labeled data is necessary, autoencoders are an unsupervised learning technique
- Autoencoders trained on language data are neural language models

# Autoencoders

- An **autoencoder** is a neural network where the size of the output layer is the same size as the input layer
- The hidden layers are usually smaller
- The goal is to generalize the training data
- Since no labeled data is necessary, autoencoders are an unsupervised learning technique
- Autoencoders trained on language data are neural language models
- Autoencoders are occasionally called diabolo networks



## Tips & Tricks (discussed in class)

- Network depth
- Layer size
- Dropout
- Early stopping
- Optimizers
- Learning rate

## Software

- Most popular neural net software are based on the following:

<b>Name</b>	<b>Lang Support</b>	<b>GPU Support</b>	<b>Who</b>
<b>Theano</b>	Python	Yes	Uni Montreal
<b>TensorFlow</b>	Python, C++	Yes	Google
<b>Torch</b>	Lua	Yes	FB, Twitter, etc.
<b>DL4J</b>	Java, Scala	Yes	Skyminid.io
<b>CNTK</b>	C++	Yes	Microsoft



## Software

- Most popular neural net software are based on the following:

<b>Name</b>	<b>Lang Support</b>	<b>GPU Support</b>	<b>Who</b>
<b>Theano</b>	Python	Yes	Uni Montreal
<b>TensorFlow</b>	Python, C++	Yes	Google
<b>Torch</b>	Lua	Yes	FB, Twitter, etc.
<b>DL4J</b>	Java, Scala	Yes	Skyminid.io
<b>CNTK</b>	C++	Yes	Microsoft

- Many others: Caffe, MXNet, CNN

## Software

- Most popular neural net software are based on the following:

Name	Lang Support	GPU Support	Who
<b>Theano</b>	Python	Yes	Uni Montreal
<b>TensorFlow</b>	Python, C++	Yes	Google
<b>Torch</b>	Lua	Yes	FB, Twitter, etc.
<b>DL4J</b>	Java, Scala	Yes	Skymind.io
<b>CNTK</b>	C++	Yes	Microsoft

- Many others: Caffe, MXNet, CNN
- We'll use Keras ([keras.io](http://keras.io)), which is really easy and intuitive. It can use either Theano or TensorFlow as a backend.